

Project Name: Localization

序号	学号	专业班级	姓名	性别
1				

1. Project Introduction

内容包括：（这部分内容不要太长，讲清楚即可）

- （1） 开发环境及系统运行要求，包括所用的开发工具、开发包、开源库、系统运行要求等；

- Linux
- Ubuntu16.04
- Python3.6.1
- CUDA 9.0
- Pytorch
- Jupyter Notebook

本次实验使用 AWS 的 p2.xlarge 服务器，在 Anaconda 中配置对应的环境。

- （2） 工作分配简介，即谁要做什么事情
独立完成

2. Technical Details

内容包括：

- （1） 工程实践当中所用到的理论知识阐述

- 模型

使用在 ImageNet 上预训练好的 18 层残差网络。

layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112×3	7×7, 64, stride 2				
		3×3 max pool, stride 2				
conv2_x	56×56×3	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3_x	28×28×3	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
conv4_x	14×14×3	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7×3	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	1×1	average pool, 1000-d fc, softmax				
FLOPs		1.8×10 ⁹	3.6×10 ⁹	3.8×10 ⁹	7.6×10 ⁹	11.3×10 ⁹

Table 1. Architectures for ImageNet. Building blocks are shown in brackets (see also Fig. 5), with the numbers of blocks stacked. Down-sampling is performed by conv3_1, conv4_1, and conv5_1 with a stride of 2.

- 损失函数

$$L_{\text{loc}}(t^u, v) = \sum_{i \in \{x, y, w, h\}} \text{smooth}_{L_1}(t_i^u - v_i),$$

in which


$$\text{smooth}_{L_1}(x) = \begin{cases} 0.5x^2 & \text{if } |x| < 1 \\ |x| - 0.5 & \text{otherwise,} \end{cases}$$

- 正确率

若预测的 bounding box 与 ground truth 的 IoU 大于等于 0.75，则认为预测正确。

- 训练

使用 Adam 算法进行优化，batch size 为 32，初始学习率设置为 1e-3，每隔 5 个 epoch 将学习率降低 10 倍，训练 20 个 epochs。



(2) 具体的算法，请用文字、示意图或者是伪代码等形式进行描述（不要

贴大段的代码)

- 准备数据:

```
dataloaders = {split: torch.utils.data.DataLoader(
    datasets[split], batch_size=32, shuffle=(split=='train'),
    num_workers=2, pin_memory=True) for split in
('train', 'test')}
```

- 构建模型

```
bar = progressbar.ProgressBar()
for ims, boxes, im_sizes in bar(dataloaders[phase]):
    boxes = crop_boxes(boxes, im_sizes)
    boxes = box_transform(boxes, im_sizes)

    inputs = Variable(ims.cuda())
    targets = Variable(boxes.cuda())

    optimizer.zero_grad()

    # forward
    outputs = model(inputs)
    loss = criterion(outputs, targets)
    acc = compute_acc(outputs.data.cpu(), targets.data.cpu(),
im_sizes)

    nsample = inputs.size(0)
    accs.update(acc, nsample)
    losses.update(loss.data[0], nsample)

    if phase == 'train':
        loss.backward()
        optimizer.step()
```

- 计算 Acc:

```
def compute_acc(preds, targets, im_sizes, theta=0.75):
    preds = box_transform_inv(preds.clone(), im_sizes)
    preds = crop_boxes(preds, im_sizes)
    targets = box_transform_inv(targets.clone(), im_sizes)
```

```
IoU = compute_IoU(preds, targets)
corr = (IoU >= theta).sum()
return corr.item() / preds.size(0)
```

- (3) 程序开发中重要的技术细节，比如用到了哪些重要的函数？这些函数来自于哪些基本库？功能是什么？自己编写了哪些重要的功能函数？等等

- torch 本次实验最重要的库，本次程序基于 pytorch 完成。

3. Experiment Results

用图文并茂的形式给出实验结果，如系统界面、操作说明、运行结果等，并对实验结果进行总结和说明。

1. 服务器：

本次实验在 AWS 端选择 p2.xlarge 类服务器。

筛选条件: GPU 计算 最新一代 显示/隐藏列

当前选择的实例类型: p2.xlarge (11.75 ECU, 4 vCPU, 2.7 GHz, E5-2686v4, 61 GiB 内存, 仅限于 EBS)


	系列	类型	vCPU	内存 (GiB)	实例存储 (GB)	可用的优化 EBS	网络性能	IPv6 支持
	GPU 计算	p2.xlarge	4	61	仅限于 EBS	是	高	是

2. Localization 结果：

```
In [20]: ind = random.choice(range(len(datasets['train'])))
im, box, im_size = datasets['train'][ind]
box = box_transform(box, im_size)

inp = im.numpy().transpose((1, 2, 0))
mean = np.array([0.485, 0.456, 0.406])
std = np.array([0.229, 0.224, 0.225])
inp = std * inp + mean
inp = np.clip(inp, 0, 1)

imshow(inp, box[0])
```



在 Jupyter Notebook 运行源代码，显示最终的 localization 结果。此结果显示训练效果较好。

3. 损失函数：

N/A% (0 of 295) | Elapsed Time: 0:00:00 ETA: --:--:--/home/ubuntu/anaconda3/envs/ai_3_2
/lib/python3.6/site-packages/ipykernel_launcher.py:51: UserWarning: invalid index of a 0-dim tensor. This will be an
error in PyTorch 0.5. Use tensor.item() to convert a 0-dim tensor to a Python number
100% (295 of 295) |#####| Elapsed Time: 0:01:17 Time: 0:01:17

[train] Epoch: 1/20 Loss: 0.0266 Acc: 24.81% Time: 78.324

100% (75 of 75) |#####| Elapsed Time: 0:00:11 Time: 0:00:11

[test] Epoch: 1/20 Loss: 0.0233 Acc: 29.53% Time: 11.691
[Info] best test acc: 29.53% at 0th epoch

100% (295 of 295) |#####| Elapsed Time: 0:01:18 Time: 0:01:18

[train] Epoch: 2/20 Loss: 0.0110 Acc: 40.33% Time: 78.663

100% (75 of 75) |#####| Elapsed Time: 0:00:11 Time: 0:00:11

[test] Epoch: 2/20 Loss: 0.0113 Acc: 37.67% Time: 11.459
[Info] best test acc: 37.67% at 1th epoch

[test] Epoch: 2/20 Loss: 0.0113 Acc: 37.67% Time: 11.459
[Info] best test acc: 37.67% at 1th epoch

100% (295 of 295) |#####| Elapsed Time: 0:01:18 Time: 0:01:18

[train] Epoch: 3/20 Loss: 0.0080 Acc: 48.92% Time: 78.759

100% (75 of 75) |#####| Elapsed Time: 0:00:11 Time: 0:00:11

[test] Epoch: 3/20 Loss: 0.0078 Acc: 52.43% Time: 11.353
[Info] best test acc: 52.43% at 2th epoch

100% (295 of 295) |#####| Elapsed Time: 0:01:18 Time: 0:01:18

[train] Epoch: 4/20 Loss: 0.0053 Acc: 63.70% Time: 78.863

100% (75 of 75) |#####| Elapsed Time: 0:00:11 Time: 0:00:11

[test] Epoch: 4/20 Loss: 0.0070 Acc: 57.50% Time: 11.732
[Info] best test acc: 57.50% at 3th epoch

100% (295 of 295) |#####| Elapsed Time: 0:01:18 Time: 0:01:18

[train] Epoch: 5/20 Loss: 0.0048 Acc: 65.62% Time: 78.820

100% (75 of 75) |#####| Elapsed Time: 0:00:11 Time: 0:00:11

[test] Epoch: 5/20 Loss: 0.0053 Acc: 66.55% Time: 11.650
[Info] best test acc: 66.55% at 4th epoch

100% (295 of 295) |#####| Elapsed Time: 0:01:18 Time: 0:01:18

[train] Epoch: 6/20 Loss: 0.0020 Acc: 90.76% Time: 78.817

100% (75 of 75) |#####| Elapsed Time: 0:00:11 Time: 0:00:11

[test] Epoch: 6/20 Loss: 0.0024 Acc: 89.67% Time: 11.846
[Info] best test acc: 89.67% at 5th epoch

100% (295 of 295) |#####| Elapsed Time: 0:01:18 Time: 0:01:18

[train] Epoch: 7/20 Loss: 0.0015 Acc: 93.00% Time: 78.786

[Info] best test acc: 93.00% at 6th epoch

100% (295 of 295) |#####| Elapsed Time: 0:01:18 Time: 0:01:18

[train] Epoch: 7/20 Loss: 0.0015 Acc: 93.00% Time: 78.786

100% (75 of 75) |#####| Elapsed Time: 0:00:11 Time: 0:00:11

[test] Epoch: 7/20 Loss: 0.0020 Acc: 91.61% Time: 12.358
[Info] best test acc: 91.61% at 6th epoch

100% (295 of 295) |#####| Elapsed Time: 0:01:18 Time: 0:01:18

[train] Epoch: 8/20 Loss: 0.0014 Acc: 94.10% Time: 78.955

100% (75 of 75) |#####| Elapsed Time: 0:00:11 Time: 0:00:11

[test] Epoch: 8/20 Loss: 0.0019 Acc: 92.54% Time: 11.633
[Info] best test acc: 92.54% at 7th epoch

100% (295 of 295) |#####| Elapsed Time: 0:01:18 Time: 0:01:18

[train] Epoch: 8/20 Loss: 0.0014 Acc: 94.88% Time: 78.962

[Info] best test acc: 92.54% at 7th epoch

100% (295 of 295) |#####| Elapsed Time: 0:01:18 Time: 0:01:18

[train] Epoch: 9/20 Loss: 0.0012 Acc: 94.88% Time: 78.962

100% (75 of 75) |#####| Elapsed Time: 0:00:10 Time: 0:00:10

[test] Epoch: 9/20 Loss: 0.0017 Acc: 93.07% Time: 11.081
[Info] best test acc: 93.07% at 8th epoch

100% (295 of 295) |#####| Elapsed Time: 0:01:18 Time: 0:01:18

[train] Epoch: 10/20 Loss: 0.0011 Acc: 95.93% Time: 78.825

100% (75 of 75) |#####| Elapsed Time: 0:00:11 Time: 0:00:11

[test] Epoch: 10/20 Loss: 0.0016 Acc: 94.10% Time: 11.683
[Info] best test acc: 94.10% at 9th epoch

100% (295 of 295) |#####| Elapsed Time: 0:01:18 Time: 0:01:18

	[Info] best test acc: 94.10% at 9th epoch
100% (295 of 295) #####	Elapsed Time: 0:01:18 Time: 0:01:18
[train] Epoch: 11/20	Loss: 0.0008 Acc: 97.20% Time: 78.832
100% (75 of 75) #####	Elapsed Time: 0:00:11 Time: 0:00:11
[test] Epoch: 11/20	Loss: 0.0014 Acc: 95.12% Time: 11.504
[Info] best test acc: 95.12% at 10th epoch	
100% (295 of 295) #####	Elapsed Time: 0:01:18 Time: 0:01:18
[train] Epoch: 12/20	Loss: 0.0008 Acc: 97.75% Time: 78.743
100% (75 of 75) #####	Elapsed Time: 0:00:11 Time: 0:00:11
[test] Epoch: 12/20	Loss: 0.0014 Acc: 95.55% Time: 11.365
[Info] best test acc: 95.55% at 11th epoch	
100% (295 of 295) #####	Elapsed Time: 0:01:18 Time: 0:01:18
[train] Epoch: 13/20	Loss: 0.0007 Acc: 97.92% Time: 78.769
100% (75 of 75) #####	Elapsed Time: 0:00:11 Time: 0:00:11
[test] Epoch: 13/20	Loss: 0.0013 Acc: 95.68% Time: 11.695
[Info] best test acc: 95.68% at 12th epoch	
100% (295 of 295) #####	Elapsed Time: 0:01:18 Time: 0:01:18
[train] Epoch: 14/20	Loss: 0.0008 Acc: 97.65% Time: 78.845
100% (75 of 75) #####	Elapsed Time: 0:00:10 Time: 0:00:10
[test] Epoch: 14/20	Loss: 0.0013 Acc: 95.49% Time: 11.242
[Info] best test acc: 95.68% at 12th epoch	
100% (295 of 295) #####	Elapsed Time: 0:01:18 Time: 0:01:18
[train] Epoch: 15/20	Loss: 0.0007 Acc: 97.49% Time: 78.840
100% (75 of 75) #####	Elapsed Time: 0:00:11 Time: 0:00:11
100% (75 of 75) #####	Elapsed Time: 0:00:11 Time: 0:00:11
[test] Epoch: 15/20	Loss: 0.0013 Acc: 95.33% Time: 11.344
[Info] best test acc: 95.68% at 12th epoch	
100% (295 of 295) #####	Elapsed Time: 0:01:18 Time: 0:01:18
[train] Epoch: 16/20	Loss: 0.0007 Acc: 98.48% Time: 78.834
100% (75 of 75) #####	Elapsed Time: 0:00:11 Time: 0:00:11
[test] Epoch: 16/20	Loss: 0.0013 Acc: 96.16% Time: 12.097
[Info] best test acc: 96.16% at 15th epoch	
100% (295 of 295) #####	Elapsed Time: 0:01:18 Time: 0:01:18
[train] Epoch: 17/20	Loss: 0.0007 Acc: 98.23% Time: 78.828
100% (75 of 75) #####	Elapsed Time: 0:00:11 Time: 0:00:11
[test] Epoch: 17/20	Loss: 0.0013 Acc: 95.99% Time: 11.516
[Info] best test acc: 96.16% at 15th epoch	
100% (75 of 75) #####	Elapsed Time: 0:00:10 Time: 0:00:10
[test] Epoch: 18/20	Loss: 0.0013 Acc: 95.78% Time: 11.275
[Info] best test acc: 96.16% at 15th epoch	
100% (295 of 295) #####	Elapsed Time: 0:01:18 Time: 0:01:18
[train] Epoch: 19/20	Loss: 0.0007 Acc: 98.09% Time: 78.796
100% (75 of 75) #####	Elapsed Time: 0:00:10 Time: 0:00:10
[test] Epoch: 19/20	Loss: 0.0013 Acc: 95.87% Time: 11.286
[Info] best test acc: 96.16% at 15th epoch	
100% (295 of 295) #####	Elapsed Time: 0:01:18 Time: 0:01:18
[train] Epoch: 20/20	Loss: 0.0007 Acc: 98.14% Time: 78.884
100% (75 of 75) #####	Elapsed Time: 0:00:11 Time: 0:00:11
[test] Epoch: 20/20	Loss: 0.0013 Acc: 95.89% Time: 11.914
[Info] best test acc: 96.16% at 15th epoch	

由上述实验结果可知，总体来说随着训练次数的增加 Acc 在增加，而 Loss 在减少。此实验现象说明，实验训练效果随着训练次数的增加而提高。但随着训练次数的不断增加，将出现过拟合结果，即虽然 training 的效果提高，但是 test 的结果反而下降。

本次实验由于时间比较赶，因此没有在源码的基础上进行优化。前期较大的经历用来配置程序运行所需要的环境。虽然对本次实验的原理没有非常清晰的了解，但是在实验过程中，学会了配置 GPU 以及远程的 Jupyter Notebook 环境，最终

可以运行本次实验原程序。

References:

给出主要的参考文献，可以是论文、网站、书籍、别人的技术报告等。

[1]: <https://arxiv.org/abs/1512.03385> 'Deep Residual Learning for Image Recognition'

备注:

代码中请给出较为详细的注释，此报告中切勿粘贴大量代码，否则扣分。